



60-DAY INTERNATIONAL AI/ML JOB SEARCH MASTER ROADMAP

For Senior Data Scientists / GenAI Engineers | India → Dubai, UK, Germany, France, Europe

Starting: 10th May 2026 | IIT Bombay M.Tech | 5 Years MNC Experience

SECTION 1 — GLOBAL STRATEGY

1.1 Country-by-Country Analysis for Indian Candidates

AE DUBAI / UAE

Factor	Details
Ease for Indians	★★★★★ — Highest. Largest Indian expat community, English-first workplace
Visa	Employment Visa sponsored by employer. No separate work permit process. 2–4 weeks
Salary Range	AED 25,000–60,000/month (USD 6,800–16,300). Tax-free
Sponsorship Likelihood	Very High — 85%+ of roles sponsor
Time to Offer	3–8 weeks typical
Best Companies	PwC, Deloitte, KPMG, Accenture, G42, Presight, Noon, Talabat, Majid Al Futtaim, ADNOC, Emirates Group, Careem
Hot Roles	GenAI Engineer, AI Product Manager, Data Scientist, ML Engineer
Remote	Rare. Relocation strongly preferred
Language	English sufficient. Arabic is a bonus but not required

GB UNITED KINGDOM

Factor	Details
Ease for Indians	★★★★ — Good. Skilled Worker Visa is well-defined
Visa	Skilled Worker Visa. Employer must be Home Office approved sponsor. 3–8 weeks processing
Salary Range	£60,000–£120,000/year. Senior GenAI: £90K–£130K
Sponsorship Likelihood	Medium-High — 60–70% of AI/ML roles willing to sponsor
Time to Offer	6–12 weeks
Best Companies	DeepMind, Wayve, Monzo, Revolut, Darktrace, Arm, Barclays, HSBC, Sky, BBC, BT, Experian
Hot Roles	ML Engineer, LLM Engineer, AI Research Engineer, NLP Engineer
Remote	Hybrid common. Some UK-based remote roles exist
Language	English only

DE GERMANY

Factor	Details
Ease for Indians	☆☆☆ — Moderate. EU Blue Card path is good but language barrier
Visa	EU Blue Card (salary threshold: ~€56,400 for AI roles). Or Job Seeker Visa (6 months to look for a job in Germany)
Salary Range	€70,000–€130,000/year gross. ~€50K–€85K net after taxes
Sponsorship Likelihood	Medium — 50–60%. Startup ecosystem growing
Time to Offer	8–16 weeks
Best Companies	SAP, Siemens, BMW, Aleph Alpha, Celonis, DeepL, N26, Zalando, HelloFresh, Delivery Hero
Hot Roles	ML Engineer, Data Scientist, AI Engineer, Research Scientist
Remote	Growing. Many German startups offer remote-first
Language	B1 German increasingly preferred but English-first tech roles exist in Berlin/Munich

FR FRANCE

Factor	Details
Ease for Indians	☆☆☆ — Moderate. Talent Passport visa is excellent for AI roles
Visa	Talent Passport (Passeport Talent) — specifically designed for highly-skilled tech workers
Salary Range	€55,000–€110,000/year gross
Sponsorship Likelihood	Medium — ~50%
Time to Offer	8–14 weeks
Best Companies	Mistral AI, Criteo, BlaBlaCar, Doctolib, Contentsquare, Hugging Face, Thales, Atos, Capgemini
Hot Roles	AI Researcher, ML Engineer, GenAI Engineer (Mistral ecosystem is booming)
Language	French preferred but major tech companies are English-friendly

NL NETHERLANDS

Factor	Details
Ease for Indians	☆☆☆☆ — Good. 30% tax ruling is a major financial benefit
Visa	Highly Skilled Migrant (HSM) visa. Fast processing (2 weeks if employer is IND-recognized)
Salary Range	€70,000–€120,000/year. 30% tax ruling = ~30% tax-free for 5 years
Sponsorship Likelihood	High — 70%+
Time to Offer	6–10 weeks
Best Companies	ASML, Philips, Booking.com, bol.com, ING, Adyen, TomTom, Elastic
Hot Roles	ML Engineer, Data Scientist, AI/GenAI Engineer
Language	English widely used in tech. Dutch not required

IE IRELAND

Factor	Details
Ease for Indians	★★★★ — Very good. Critical Skills Employment Permit fast-tracked
Visa	Critical Skills Employment Permit (€38,000+ salary threshold for eligible occupations)
Salary Range	€70,000–€120,000/year
Sponsorship Likelihood	High — 75%+
Best Companies	Google, Meta, Microsoft, LinkedIn, Stripe, HubSpot, Zalando Dublin, Accenture Dublin
Hot Roles	ML Engineer, AI Engineer, Data Scientist (at big tech EU hubs)

1.2 Salary Expectations Summary Table

Country	Junior DS (3y)	Senior DS (5y)	GenAI/LLM Eng	Tax Burden	Take-Home Est.
Dubai/UAE	AED 18–25K/mo	AED 28–50K/mo	AED 35–60K/mo	0%	Full salary
UK	£55–70K/yr	£75–100K/yr	£90–130K/yr	~35–40%	£45–75K/yr
Germany	€60–80K/yr	€80–110K/yr	€90–130K/yr	~35–42%	€45–70K/yr
Netherlands	€65–85K/yr	€85–110K/yr	€95–130K/yr	~25–35%*	€55–85K/yr
France	€55–75K/yr	€75–100K/yr	€85–120K/yr	~30–40%	€45–65K/yr
Ireland	€65–85K/yr	€85–110K/yr	€90–125K/yr	~35–40%	€50–70K/yr

*Netherlands 30% ruling applies if eligible

Conclusion: Dubai offers highest take-home. Netherlands with 30% ruling is best net in Europe.

1.3 Target Role Priority Matrix

Role Title	UAE	UK	Germany	Demand 2026	Your Fit
GenAI / LLM Engineer	★★★★★	★★★★★	★★★★★	Very High	✔ Perfect
AI Engineer (Agentic)	★★★★★	★★★★★	★★★★★	Very High	✔ Perfect
Senior Data Scientist	★★★★	★★★★★	★★★★★	High	✔ Strong
ML Engineer	★★★	★★★★★	★★★★★	Very High	✔ Good
AI Research Scientist	★★	★★★★★	★★★★★	High	⚠ PhD preferred
MLOps Engineer	★★★	★★★★★	★★★★★	High	✔ Good

Primary title to use: "Senior AI/LLM Engineer" or "GenAI Engineer" — highest demand + best match

1.4 Company Type Strategy

Startups (Recommended: 40% of applications)

Why: Move fast, value demonstrated skills over pedigree, genuinely sponsor visas, equity upside Target: Series B–D AI startups in Dubai, London, Amsterdam, Berlin Examples: Wayve, Cohere, Poolside, ElevenLabs, Mistral (Paris), Convergence AI

Consulting/Big 4 (Recommended: 30% of applications)

Why: Structured relocation packages, volume hiring, visa expertise **Target:** Accenture AI, McKinsey QuantumBlack, BCG X, Deloitte AI, PwC AI **Advantage:** Your MNC background aligns perfectly

Product Companies/Scaleups (Recommended: 20% of applications)

Why: Best salary + equity, interesting problems **Target:** Noon, Careem (UAE), Monzo/Revolut (UK), Zalando/Delivery Hero (Germany)

Big Tech (Recommended: 10% — long shot but worth trying)

Why: Highest prestige, strong relocation packages **Target:** Google, Meta, Microsoft EU offices (Dublin/Amsterdam/London)

1.5 2026 Hiring Trends in AI/ML

What companies are actively hiring for:

1. **Agentic AI / Multi-Agent Systems** — LangGraph, AutoGen, CrewAI patterns
2. **Production RAG pipelines** — Not just POCs, but enterprise-grade RAG with eval frameworks
3. **LLM Fine-tuning** — QLoRA, LoRA, PEFT on open-source models (Llama, Mistral)
4. **AI Infrastructure** — Model serving, vLLM, TGI, batching strategies
5. **AI Evaluation / Evals** — LLM-as-judge, RAGAS, automated evaluation pipelines
6. **Multimodal AI** — Vision-language models, document AI
7. **AI Safety / Guardrails** — NeMo Guardrails, constitutional AI, hallucination reduction

Skills in decline (avoid overemphasizing):

- Traditional ML only (scikit-learn pipelines without LLM context)
 - Jupyter-notebook data analysis as primary work
 - BI/visualization-heavy roles
-

1.6 Resume Strategy by Country

UAE Resume

- **Length:** 2 pages max
- **Photo:** Not required but not penalized
- **DOB/Marital:** Not needed
- **Style:** Achievement-focused, quantified metrics, clean formatting
- **Nationality:** Many job postings ask; include "Indian National / Open to UAE Visa"
- **Notice Period:** Mention if ≤30 days or "Immediate"
- **ATS:** Dubai companies often manually review, but still ATS-optimize

UK Resume (Called CV)

- **Length:** 2 pages strict
- **Photo:** Do NOT include
- **DOB/Marital:** Do NOT include (legal requirement)
- **Style:** Reverse chronological, concise, no "objective statement"

- **Visa:** Add "Requires Skilled Worker Visa Sponsorship" in header
- **ATS:** Very ATS-heavy. Use exact keywords from JD

Germany Resume (Lebenslauf)

- **Length:** 2 pages + cover letter (mandatory)
- **Photo:** Professional photo is expected and traditional
- **DOB:** Often included (traditional expectation)
- **Style:** Very structured, formal, table-like format acceptable
- **Cover Letter:** 1 page, tailored, in English (unless specified German)
- **Certificates:** Scan and attach degree/transcripts if requested

France Resume (CV)

- **Length:** 1–2 pages
- **Photo:** Common/expected
- **Style:** Clean, minimal. French appreciate elegant formatting
- **Language:** English is fine for international companies; French CV for local firms
- **Motivation Letter:** Expected; brief and personal

1.7 LinkedIn Optimization Strategy

Profile Score Target: All-Star

1. **Headline** (Most Critical):
 -  Bad: "Data Scientist at XYZ"
 -  Good: "Senior AI/LLM Engineer | GenAI • Multi-Agent Systems • RAG Pipelines | Open to Dubai/UK/Europe"
2. **About Section:** 3–5 paragraphs. Start with impact, not backstory. Include: what you build, tech stack, scale achieved, what you're seeking.
3. **Open to Work** (Recruiter-only mode, NOT public banner):
 - Job titles: "GenAI Engineer", "LLM Engineer", "Senior Data Scientist", "AI Engineer"
 - Locations: Dubai, London, Amsterdam, Berlin, Paris, Remote
 - Start date: Immediately / Within 1 month
4. **Experience:** Every role needs 3–5 bullets with metrics. "Built RAG pipeline serving 10K daily queries" > "Developed RAG pipeline"
5. **Skills:** Add ALL relevant keywords. Top skills to list:
 - Large Language Models, Retrieval-Augmented Generation, LangChain, LangGraph, Python, FastAPI, GCP, Vertex AI, Vector Databases, MLOps, Deep Learning, NLP, PyTorch, Transformers
6. **Featured Section:** Pin GitHub portfolio link, any demo video, any publication
7. **Activity:** Post 2–3 times/week on LinkedIn. Technical posts about GenAI topics get massive reach.

1.8 Cold Outreach & Recruiter Strategy

Recruiter DM Template (LinkedIn)



Subject: Senior AI/LLM Engineer | 5 YOE | Open to [Country] Relocation

Hi [Name],

I came across your profile while researching AI/ML hiring at [Company].

Quick background: I'm a Senior AI/LLM Engineer with 5 years at [MNC], IIT Bombay M.Tech, specializing in production GenAI systems — RAG pipelines, multi-agent orchestration (LangGraph), Vertex AI, FastAPI backends. Currently open to relocating to [Dubai/UK/Germany].

Would love to explore if there are any relevant openings on your team. Happy to share my resume or have a quick call.

Best,
[Name]

Target Recruiters to Connect With:

- **UAE:** Hays UAE, Robert Half Dubai, Michael Page Gulf, Antal International
- **UK:** Hays Technology, Reed Technology, Nigel Frank, Real Staffing
- **Germany:** Hays Germany, Michael Page Germany, Experteer
- **Netherlands:** YER, Brunel, Michael Page Netherlands

Key LinkedIn Boolean Search to Find Recruiters:



"technical recruiter" AND ("data scientist" OR "machine learning" OR "AI") AND ("Dubai" OR "London" OR "Amsterdam")

SECTION 2 — 60-DAY DAY-WISE MASTER PLAN

Starting: 10th May 2026

WEEK 1 (May 10–16): FOUNDATION WEEK

Theme: Set up everything. Zero applications yet. Weekly Goal: 100% setup, 0% applications, 20% technical study begins

DAY 1 — Saturday, May 10

Priority: CRITICAL | Focus: LinkedIn + Resume Baseline

Time	Task	Hours
Morning (9–12)	LinkedIn audit: Update headline, About, all 5 experiences with metrics. Set Open to Work (recruiter-only)	3h
Afternoon (1–5)	Create master resume v1.0. Clean ATS format. Quantify every bullet. Export as PDF + DOCX	4h
Night (8–10)	GitHub audit: Clean repos, write READMEs for top 3 projects. Pin 3 best repos	2h
Total		9h

Applications: 0 | **Outreach:** 0 | **Study:** 0

DAY 2 — Sunday, May 11

Priority: HIGH | **Focus:** Portfolio + Application Accounts

Time	Task	Hours
Morning (9–12)	Create accounts: LinkedIn Premium (1 month free trial), Indeed, Naukri Gulf, GulfTalent, Bayt, Wellfound	2h
Afternoon (1–4)	Set up Job Application Tracker in Excel/Notion (see Section 9 template)	1h
Afternoon (4–6)	Study: Python revision — list comprehensions, generators, decorators, context managers	2h
Night (8–10)	Research: Identify 20 target companies in UAE, 20 in UK, 10 in Germany. Add to tracker	2h
Total		7h

Applications: 0 | **Outreach:** 0 | **Study:** 2h

DAY 3 — Monday, May 12

Priority: HIGH | **Focus:** Technical Study Begins + First Outreach

Time	Task	Hours
Morning (7–9)	Python: OOP deep dive — metaclasses, abstract classes, design patterns	2h
Afternoon (1–5, work hours around)	Connect with 10 recruiters in UAE on LinkedIn (use DM template)	1h
Evening (6–8)	SQL: Window functions — ROW_NUMBER, RANK, LAG, LEAD, NTILE. Practice 10 questions on LeetCode	2h
Night (9–10)	Review JDs: Collect 10 UAE job descriptions. Extract recurring keywords	1h
Total		6h

Applications: 0 | **Outreach:** 10 recruiter connections | **Study:** 4h

DAY 4 — Tuesday, May 13

Priority: HIGH | **Focus:** Statistics + Resume UAE version

Time	Task	Hours
Morning (7–9)	Statistics: Probability distributions, CLT, hypothesis testing, p-values, confidence intervals	2h
Evening (6–8)	Create UAE-specific resume v1.1 (add nationality note, optimize for Gulf companies)	2h
Night (8–10)	ML Concepts: Bias-variance tradeoff, regularization (L1/L2), cross-validation	2h
Total		6h
Applications: 0 Outreach: 10 more UK recruiter connections Study: 4h		

DAY 5 — Wednesday, May 14

Priority: HIGH | Focus: GenAI Study + LinkedIn Post #1

Time	Task	Hours
Morning (7–9)	GenAI: RAG architecture deep dive — chunking strategies, embedding models, retrieval methods	2h
Evening (6–7)	Write + publish LinkedIn post: "5 lessons from building production RAG systems"	1h
Night (8–10)	LangChain/LangGraph: Revise chains, agents, memory, tool calling. Read LangGraph docs	2h
Total		5h
Applications: 0 Outreach: 0 Study: 4h		

DAY 6 — Thursday, May 15

Priority: HIGH | Focus: System Design + Mock Interview Prep

Time	Task	Hours
Morning (7–9)	System Design basics: CAP theorem, microservices, API design, load balancing	2h
Evening (6–8)	ML System Design: Feature stores, model serving patterns, A/B testing infrastructure	2h
Night (8–9)	Behavioral prep: Write 5 STAR stories (Situation, Task, Action, Result) from your 5 years	1h
Total		5h
Applications: 0 Outreach: Follow up on Day 3 outreach Study: 4h		

DAY 7 — Friday, May 16

Priority: MEDIUM | Focus: Review Week + First 5 Applications

Time	Task	Hours
Morning (9–11)	Weekly review: Assess tracker, profile views, recruiter responses	1h
Morning (11–1)	Apply to 5 carefully selected UAE roles (not spray-and-pray — tailored applications)	2h
Afternoon (2–5)	Deep Learning revision: Backpropagation, CNN architectures, attention mechanism	3h
Night (8–10)	Vector Databases: Study Pinecone, Weaviate, Qdrant. Understand HNSW, ANN search	2h
Total		8h
Applications: 5 Outreach: 0 Study: 5h		

WEEK 1 MILESTONES

- ☐ LinkedIn All-Star status achieved
- ☐ Master resume v1.0 ready
- ☐ UAE resume v1.1 ready
- ☐ 20 recruiters connected (UAE + UK)
- ☐ Job application tracker set up
- ☐ 50 target companies identified
- ☐ 5 applications submitted
- ☐ 20h+ technical study completed
- ☐ GitHub cleaned and pinned

Week 1 KPIs:

- LinkedIn Profile Views: Target 50+ (from baseline)
- Recruiter connections accepted: Target 15/20
- Application response rate: 0% expected (too early)

WEEK 2 (May 17–23): MOMENTUM WEEK

Theme: Ramp applications, deepen technical prep, build portfolio Weekly Goal: 20+ applications, first recruiter conversations, Project 1 started

DAY 8 — Saturday, May 17

Priority: CRITICAL | Focus: Project 1 Start — Production RAG System

Time	Task	Hours
Morning (9–12)	Design architecture for Project 1: Enterprise RAG with eval framework. Write README first	3h
Afternoon (1–5)	Build: Document ingestion pipeline (PDF/DOCX), chunking, embedding (OpenAI/Vertex AI)	4h
Night (8–10)	Apply to 5 UK roles. Tailor CV for each. Track in spreadsheet	2h
Total		9h

DAY 9 — Sunday, May 18

Priority: HIGH | Focus: RAG Project + LLM Study

Time	Task	Hours
Morning (9–12)	Continue Project 1: Retrieval layer (Qdrant/Pinecone), reranking (Cohere/cross-encoder)	3h
Afternoon (1–3)	LLM Study: Tokenization, BPE, attention, transformers architecture in detail	2h
Afternoon (3–5)	LLM Study: Fine-tuning methods — full fine-tuning, LoRA, QLoRA, PEFT	2h
Night (8–10)	Apply to 5 more UAE/Gulf roles	2h
Total		9h

DAY 10 — Monday, May 19

Priority: HIGH | Focus: FastAPI + Agentic Systems

Time	Task	Hours
Morning (7–9)	FastAPI: Dependency injection, background tasks, middleware, auth (JWT), async routes	2h
Evening (6–8)	LangGraph study: State graphs, nodes, edges, conditional routing, human-in-the-loop	2h
Night (8–10)	Apply to 5 Germany/Netherlands roles. Research EU Blue Card salary requirement	2h
Total		6h

DAY 11 — Tuesday, May 20

Priority: HIGH | Focus: Vector DBs + RAG Evals

Time	Task	Hours
Morning (7–9)	Vector DB deep dive: HNSW index, product quantization, hybrid search (dense+sparse)	2h
Evening (6–8)	RAG Evaluation: RAGAS framework, faithfulness, context precision, answer relevancy	2h
Night (8–10)	Project 1: Add evaluation module using RAGAS. Push to GitHub	2h
Total		6h

DAY 12 — Wednesday, May 21

Priority: HIGH | Focus: Cold Outreach Blitz + LinkedIn Post

Time	Task	Hours
Morning (7–9)	Statistics: Bayesian inference, A/B testing, MLE, MAP estimation	2h
Evening (6–7)	LinkedIn post #2: "Why RAG evaluation matters more than RAG retrieval"	1h
Evening (7–9)	Cold outreach: Message 20 senior engineers/hiring managers at target companies in UAE/UK	2h
Night (9–10)	Coding practice: LeetCode — Arrays, Two Pointers (5 problems)	1h
Total		6h

DAY 13 — Thursday, May 22

Priority: HIGH | Focus: ML Study + Behavioral Prep

Time	Task	Hours
Morning (7–9)	ML: Tree methods — Random Forest, GBM, XGBoost, LightGBM. Feature importance	2h
Evening (6–8)	Behavioral interview prep: Add 5 more STAR stories. Focus on "Lead without authority", "Conflict resolution", "Ambiguous projects"	2h
Night (8–10)	Apply to 5 Ireland/France roles	2h
Total		6h

DAY 14 — Friday, May 23

Priority: MEDIUM | Focus: Week 2 Review + Mock Interview

Time	Task	Hours
Morning (9–11)	Self mock interview (record yourself): 5 technical questions + 3 behavioral	2h
Afternoon (2–5)	ML System Design practice: Design a "Recommendation System for e-commerce" end-to-end	3h
Night (8–10)	Review week: Update tracker, count applications, assess response rate	1h
Total		6h

WEEK 2 MILESTONES

- ☐ 20+ total applications submitted
- ☐ Project 1 (RAG) 60% complete
- ☐ 2 LinkedIn posts published
- ☐ 20 cold outreach messages sent
- ☐ LLM architecture fully revised
- ☐ 5 STAR stories documented
- ☐ First mock interview done

WEEK 3 (May 24–30): ACCELERATION WEEK

Theme: More applications, project completion, recruiter calls begin Weekly Goal: 35+ total applications, first recruiter calls, Project 2 start

DAY 15 — Saturday, May 24

Priority: CRITICAL | Focus: Project 1 Finalize + Project 2 Start

Time	Task	Hours
Morning (9–12)	Finalize Project 1: FastAPI backend, Streamlit frontend, Docker, deploy to GCP Cloud Run	3h
Afternoon (1–4)	Start Project 2: Multi-Agent AI System. Design architecture: Orchestrator + 3 specialized agents	3h
Night (8–10)	Apply 5 applications. Follow up on all week 1-2 applications (email/LinkedIn)	2h
Total		8h

DAY 16 — Sunday, May 25

Priority: HIGH | Focus: Deep Learning + NLP

Time	Task	Hours
Morning (9–12)	Deep Learning: Transformers in detail — multi-head attention, positional encoding, layer norm	3h
Afternoon (1–4)	NLP: Named entity recognition, relation extraction, semantic similarity, sentence embeddings	3h
Night (8–10)	Project 2: Build orchestrator agent with LangGraph. Tool definition for sub-agents	2h
Total		8h

DAY 17 — Monday, May 26

Priority: HIGH | Focus: Cloud + MLOps

Time	Task	Hours
Morning (7–9)	GCP deep dive: Cloud Run, Vertex AI Pipelines, BigQuery ML, Pub/Sub, Cloud Storage	2h
Evening (6–8)	MLOps: MLflow, model registry, model versioning, CI/CD for ML pipelines	2h
Night (8–9)	Apply 3 more roles. Research salary ranges for LinkedIn salary settings	1h
Total		5h

DAY 18 — Tuesday, May 27

Priority: HIGH | Focus: Coding Practice Intensive

Time	Task	Hours
Morning (7–9)	LeetCode: Sliding Window (5 problems), Hash Maps (5 problems)	2h
Evening (6–8)	ML Coding: Implement logistic regression from scratch, gradient descent, backprop	2h
Night (8–10)	LinkedIn post #3: "My architecture for building multi-agent AI systems in production"	1h
Total		5h

DAY 19 — Wednesday, May 28

Priority: CRITICAL | Focus: Recruiter Call Prep

Time	Task	Hours
Morning (7–9)	Prepare recruiter call script: elevator pitch (60 sec), key highlights, salary expectation, visa discussion	2h
Evening (6–8)	Study: Docker fundamentals — Dockerfile, docker-compose, multi-stage builds, container networking	2h
Night (8–10)	Research: For each target company, know: products, recent AI news, tech stack, recent funding	2h
Total		6h

DAY 20 — Thursday, May 29

Priority: HIGH | Focus: System Design + Architecture

Time	Task	Hours
Morning (7–9)	ML System Design: Design "Real-time fraud detection system" — feature engineering, model serving, monitoring	2h
Evening (6–8)	System Design: Design scalable REST API with caching, rate limiting, auth	2h
Night (8–9)	Apply 3 more roles	1h
Total		5h

DAY 21 — Friday, May 30

Priority: MEDIUM | Focus: Week 3 Review + Mock Technical Interview

Time	Task	Hours
Morning (9–12)	Mock technical interview (use Pramp, Interviewing.io, or ask friend): ML concepts + 1 coding problem	3h
Afternoon (2–4)	Weekly review: response rate analysis, update resume if needed	2h
Night (8–10)	Catch up on any applications	2h
Total		7h

WEEK 3 MILESTONES

- ☐ 35+ total applications
- ☐ Project 1 (RAG) 100% complete + deployed on GCP
- ☐ Project 2 (Multi-Agent) 50% complete
- ☐ 3 LinkedIn posts, 500+ impressions each
- ☐ Recruiter call script ready
- ☐ First recruiter call likely (if Week 1-2 outreach worked)
- ☐ LeetCode: 20+ problems solved

WEEK 4 (May 31 – June 6): INTERVIEW PREP INTENSIVE

Theme: First interviews may start. Deep technical prep. Weekly Goal: 50+ total applications, 2–3 recruiter screens expected

DAY 22–23: May 31–June 1 (Weekend)

Focus: LLM Fine-tuning + Project 2 Completion

Task	Hours
LLM Fine-tuning workshop: QLoRA hands-on with Llama/Mistral on Google Colab	6h
Complete Project 2: Multi-agent system with LangGraph, deploy API	4h
Apply 10 roles across all target countries	2h
LinkedIn post: "How I fine-tuned Llama 3 on custom data with QLoRA"	1h

DAY 24–26: June 2–4 (Weekdays)

Focus: Technical Interview Prep — ML Concepts Each day: 2h morning study + 1h LeetCode + 1h apply

Day	Topic
June 2	Embeddings deep dive: word2vec, GloVe, BERT embeddings, contrastive learning
June 3	Prompt engineering: zero-shot, few-shot, chain-of-thought, structured outputs
June 4	AI evaluation: BLEU, ROUGE, METEOR, LLM-as-judge, human eval frameworks

DAY 27–28: June 5–6

Focus: Full Mock Interview Day (Saturday) + Project 3 Start (Sunday)

- Saturday: Full mock interview (HR screen + technical round + system design) — 4h
- Sunday: Start Project 3 (AI Automation / n8n + LLM workflow) — 4h
- Apply 5 more roles

WEEK 4 MILESTONES

- ☐ 50+ total applications across all countries
 - ☐ Project 2 deployed
 - ☐ First real recruiter screen completed
 - ☐ Full mock interview done
 - ☐ 30+ LeetCode problems solved
-

WEEK 5 (June 7–13): FIRST INTERVIEWS WEEK

Theme: Real technical interviews begin. Project 3 completes. **Weekly Goal:** 2–3 technical interview rounds

Daily Schedule (Week 5 Weekdays):

Time	Activity
7:00–9:00 AM	Deep technical study (LLM/RAG/Agents topic)
9:00–6:00 PM	Day job
6:30–8:30 PM	Interview prep specific to scheduled interviews
8:30–9:30 PM	Apply 3 roles + follow ups
9:30–10:30 PM	LeetCode (2 problems)

Study Topics Week 5:

- Monday: Kubernetes basics — pods, deployments, services, ingress, HPA
- Tuesday: Distributed systems — consistency models, sharding, message queues
- Wednesday: Model monitoring — data drift, concept drift, monitoring frameworks
- Thursday: AI security — prompt injection, adversarial examples, model extraction
- Friday: Salary negotiation tactics — anchoring, competing offers, visa negotiation

WEEK 5 MILESTONES

- ☐ 65+ total applications
 - ☐ Project 3 complete
 - ☐ 2–3 technical interviews done
 - ☐ 40+ LeetCode problems
 - ☐ Salary negotiation script ready
-

WEEK 6 (June 14–20): DEEP INTERVIEW PREP

Theme: Convert interviews to offers. Negotiate.

DAY 36–42 Focus Areas:

Day	Primary Focus	Secondary
Day 36	ML System Design: Design "LLM-powered customer support" end-to-end	Apply 5 roles
Day 37	Coding: Graphs + Dynamic Programming (LeetCode medium)	Mock behavioral
Day 38	GenAI case study prep: Practice "Design RAG for enterprise" case	Apply 5 roles
Day 39	Project 4 start: Production ML system with MLflow + Feast + Seldon	Follow ups
Day 40	Data Engineering basics: dbt, Airflow, Kafka, data modeling	Apply 5 roles
Day 41	Full mock interview #3 (all rounds simulated)	Resume v3.0
Day 42	Week review + rest	LinkedIn post

WEEK 6 MILESTONES

- ☐ 80+ total applications
 - ☐ 5+ total interviews/screens completed
 - ☐ First offer potential (UAE is fastest — 4–8 week cycle)
 - ☐ Project 4 started
-

WEEKS 7–9 (June 21 – July 10): OFFER CONVERSION

Theme: Close offers, negotiate salaries, manage multiple processes

Week 7–9 Key Activities:

- Apply 5–7 roles/day (maintain pipeline volume)
- Prep deeply for each specific company's interview style
- Begin visa/relocation research
- Negotiate offers aggressively (see Section 7)
- Complete Projects 4 and 5

WEEK 7 MILESTONES

- ☐ 100+ total applications
 - ☐ 8–10 total interview rounds
 - ☐ 1–2 offers potentially in hand (UAE typical cycle)
 - ☐ GitHub: 4 production projects deployed
-

MASTER 60-DAY TRACKER

Week	Applications	Outreach	Recruiter Calls	Tech Interviews	Offers
Week 1 (May 10–16)	5	20	0	0	0
Week 2 (May 17–23)	15	20	0–1	0	0
Week 3 (May 24–30)	15	15	1–2	0–1	0
Week 4 (May 31–Jun 6)	15	10	2–3	1–2	0
Week 5 (Jun 7–13)	15	10	2–3	2–3	0
Week 6 (Jun 14–20)	15	10	1–2	2–4	0–1
Week 7 (Jun 21–27)	15	5	1–2	2–4	0–2
Week 8 (Jun 28–Jul 4)	10	5	1–2	2–4	1–2
Week 9 (Jul 5–10)	10	5	1–2	1–3	1–3
TOTAL	115+	100	10–17	10–21	1–3+

DAILY CHECKLIST (Every Weekday)



- ☐ 7–9 AM: Deep study (2h)
- ☐ 6–8 PM: Interview prep / project work (2h)
- ☐ 8–9 PM: Apply 2–3 roles (1h)
- ☐ 9–10 PM: LeetCode 2 problems OR recruiter outreach (1h)
- ☐ Log all activity in tracker
- ☐ Check LinkedIn messages and respond

DAILY CHECKLIST (Every Weekend Day)



- ☐ 9 AM–12 PM: Project work OR intensive study (3h)
- ☐ 1–5 PM: Project work OR deep technical prep (4h)
- ☐ 5–7 PM: Apply 5 roles with tailored cover notes (2h)
- ☐ 7–8 PM: LinkedIn post / content creation
- ☐ 8–10 PM: Mock interview OR system design practice (2h)

SECTION 3 — TECHNICAL PREPARATION ROADMAP

Complete Interview Prep by Topic

3.1 Python

Factor	Details
Importance	★★★★★
Interview Frequency	Every single interview
Study Hours	15–20h

Topics to Master:

- Data structures: list, dict, set, defaultdict, Counter, deque
- OOP: classes, inheritance, magic methods, metaclasses
- Functional: map, filter, reduce, lambda, generators, itertools
- Concurrency: threading vs multiprocessing vs asyncio
- Decorators, context managers, type hints
- Memory management, garbage collection, profiling
- Testing: pytest, mocking, test fixtures

Common Interview Questions:

1. "What's the GIL and how does it affect multiprocessing vs multithreading?"
2. "Implement a LRU cache using Python"
3. "What's the difference between `__str__` and `__repr__`?"
4. "How does Python manage memory? Explain reference counting + cyclic garbage collector"
5. "Write a generator that yields Fibonacci numbers indefinitely"

Best Resources:

- Fluent Python (book) — Chapters 1–11
- Real Python (realpython.com) — advanced topics
- GitHub: dabeaz/python-cookbook

3.2 SQL

Factor	Details
Importance	★★★★★
Interview Frequency	80% of data roles
Study Hours	10–15h

Topics to Master:

- Window functions: ROW_NUMBER, RANK, DENSE_RANK, LAG, LEAD, NTILE, FIRST_VALUE
- CTEs (WITH clauses), recursive CTEs
- Subqueries: correlated vs non-correlated
- Aggregations: GROUP BY, HAVING, ROLLUP, CUBE
- JOINS: INNER, LEFT, RIGHT, FULL OUTER, CROSS, SELF
- EXPLAIN/query optimization, index types (B-tree, hash)
- Transactions, ACID properties, isolation levels

Common Interview Questions:

1. "Find the second highest salary in each department"
2. "Calculate 7-day rolling average of daily sales"
3. "Find all employees who earn more than their manager"

4. "Identify sessions from clickstream data (gap > 30 min = new session)"

Best Resources:

- Mode Analytics SQL Tutorial
 - LeetCode SQL (top 50)
 - `select * from newsletter`
-

3.3 Statistics & Probability

Factor	Details
Importance	★★★★★
Interview Frequency	90% of DS roles
Study Hours	20–25h

Topics to Master:

- Distributions: Normal, Binomial, Poisson, Exponential, Beta
- Hypothesis testing: t-test, chi-square, ANOVA, Mann-Whitney U
- p-values, Type I/II errors, statistical power, effect size
- Bayesian inference: Bayes theorem, prior/posterior, MCMC
- A/B testing: sample size calculation, sequential testing, multiple comparisons
- Regression: OLS assumptions, heteroscedasticity, multicollinearity
- Causal inference: difference-in-differences, RCTs, observational studies

Common Interview Questions:

1. "Design an A/B test for a new recommendation algorithm. How do you determine sample size?"
 2. "What's the difference between Type I and Type II errors? When is each more costly?"
 3. "Explain Simpson's Paradox with a real example"
 4. "How would you handle multiple comparisons in A/B testing?"
 5. "What is the Central Limit Theorem and why does it matter practically?"
-

3.4 Machine Learning

Factor	Details
Importance	★★★★★
Interview Frequency	Every DS/ML role
Study Hours	25–30h

Core Topics:

- Linear/Logistic Regression from scratch (implement in numpy)
- Tree methods: Decision Tree, Random Forest, GBM, XGBoost, LightGBM, CatBoost
- SVM: kernel trick, soft margin, C parameter
- Clustering: K-means, DBSCAN, hierarchical, GMM
- Dimensionality reduction: PCA, t-SNE, UMAP
- Optimization: SGD, Adam, RMSprop, learning rate scheduling
- Regularization: L1 (Lasso), L2 (Ridge), ElasticNet, Dropout
- Evaluation: Precision/Recall/F1, ROC-AUC, PR-AUC, Cohen's Kappa
- Class imbalance: SMOTE, weighted loss, threshold tuning

Advanced Topics:

- Gradient boosting math (how does XGBoost differ from GBM?)
- Feature selection methods: mutual information, SHAP, permutation importance
- Hyperparameter tuning: Grid search, Bayesian optimization (Optuna)
- Model calibration: Platt scaling, isotonic regression

Best Resources:

- "Hands-On ML with Scikit-Learn & TensorFlow" (Géron)
- Stanford CS229 lecture notes
- `mlcourse.ai` open course

3.5 Deep Learning

Factor	Details
Importance	★★★★★
Interview Frequency	All GenAI roles
Study Hours	20–25h

Topics:

- Backpropagation: compute gradients manually for small networks
- CNN architectures: ResNet, EfficientNet, Vision Transformer (ViT)
- RNN/LSTM/GRU: vanishing gradients, gating mechanisms
- Attention mechanism: scaled dot-product, multi-head attention
- Transformer architecture: encoder-decoder, positional encoding, layer norm
- Batch/Layer/Group normalization: when to use which
- Loss functions: CrossEntropy, BCE, Triplet Loss, Contrastive Loss
- PyTorch: custom Dataset, DataLoader, training loops, `torch.compile`

Best Resources:

- Andrej Karpathy's "Neural Networks: Zero to Hero" (YouTube) ← MUST WATCH
- `fast.ai` Practical Deep Learning
- "Understanding Deep Learning" (Prince) — free PDF

3.6 LLM & Generative AI

Factor	Details
Importance	★★★★★
Interview Frequency	Every GenAI role
Study Hours	30–40h

Core LLM Topics:

- Pre-training: next token prediction, BERT masked LM, GPT causal LM
- RLHF: reward modeling, PPO, DPO (Direct Preference Optimization)
- Tokenization: BPE, WordPiece, SentencePiece
- Context window: position encodings, RoPE, ALiBi, sliding window attention
- Inference optimization: KV cache, speculative decoding, quantization (INT8, GPTQ, AWQ)

- Model serving: vLLM, TGI (Text Generation Inference), TensorRT-LLM
- Prompt engineering: system prompts, few-shot, chain-of-thought, ReAct

Fine-tuning:

- Full fine-tuning vs parameter-efficient (LoRA, QLoRA, IA3, Adapters)
- Instruction tuning vs RLHF vs DPO
- Data preparation: instruction templates, chat formats
- Evaluation: perplexity, MMLU, HumanEval, HellaSwag

Common Interview Questions:

1. "Explain the transformer architecture from first principles"
2. "What is LoRA? How does it reduce trainable parameters?"
3. "What's the difference between RAG and fine-tuning? When would you choose each?"
4. "Explain attention complexity. How does flash attention help?"
5. "How would you reduce LLM hallucinations in a production system?"
6. "What is DPO and how does it differ from PPO-based RLHF?"

Best Resources:

- Andrej Karpathy "Let's build GPT" (YouTube)
- Hugging Face NLP Course (huggingface.co/course)
- LLM-Engineers-Handbook (GitHub) ← Excellent
- Sebastian Raschka's "Build LLMs From Scratch" book

3.7 RAG Systems

Factor	Details
Importance	★★★★★
Interview Frequency	95% of GenAI Engineer roles
Study Hours	15–20h

Architecture Layers:

1. **Ingestion:** Document loading, parsing (Unstructured, Docling), chunking strategies
2. **Chunking:** Fixed-size, semantic, late chunking, parent-document retrieval
3. **Embedding:** Dense (OpenAI, Cohere, BGE), sparse (BM25), late interaction (ColBERT)
4. **Retrieval:** Similarity search, hybrid search (dense + sparse), MMR
5. **Reranking:** Cross-encoders (Cohere Rerank, BGE Reranker)
6. **Generation:** Context stuffing, lost-in-the-middle problem, citation generation
7. **Evaluation:** RAGAS — faithfulness, context recall, context precision, answer relevancy

Advanced RAG Patterns:

- HyDE (Hypothetical Document Embeddings)
- FLARE (Forward-Looking Active Retrieval)
- Corrective RAG (CRAG)
- Self-RAG (self-reflection on retrieval quality)
- Agentic RAG (agent decides when/what to retrieve)

Best Resources:

- LangChain RAG documentation

- "Advanced RAG" blog series (Towards Data Science)
- GitHub: `run-llama/llama_index` advanced guides

3.8 AI Agents / Multi-Agent Systems

Factor	Details
Importance	★★★★★
Interview Frequency	Growing rapidly — 70% of new GenAI roles
Study Hours	20–25h

Agent Fundamentals:

- ReAct pattern: Reason + Act loop
- Tool calling: function calling in OpenAI/Gemini/Anthropic APIs
- Memory types: in-context, external (vector store), episodic, semantic
- Planning: task decomposition, MCTS, tree-of-thought

LangGraph Specifics:

- State graphs: TypedDict state, reducers
- Nodes: regular functions and LLM nodes
- Edges: conditional edges, routing logic
- Subgraphs and nested graphs
- Human-in-the-loop: interrupt, resume
- Streaming: token-level, node-level, event streaming
- Persistence: checkpointers (SQLite, PostgreSQL)

Multi-Agent Patterns:

- Supervisor pattern (orchestrator + workers)
- Hierarchical agents
- Collaborative agents (debate, critique)
- Competitive agents
- Tool-use agents vs code-interpreting agents

Common Interview Questions:

1. "Design a multi-agent system for automated market research"
2. "How do you prevent infinite loops in an agentic system?"
3. "Explain the difference between a chain and an agent"
4. "How do you handle tool calling errors in an agent?"
5. "How would you add human approval to an AI agent workflow?"

3.9 Vector Databases

Factor	Details
Importance	★★★★★
Interview Frequency	80% of RAG/GenAI roles
Study Hours	8–10h

Understand:

- ANN algorithms: HNSW, IVF, PQ, FAISS
- Similarity metrics: cosine, dot product, Euclidean — when to use each
- Metadata filtering: pre-filter vs post-filter
- Hybrid search: combining vector search with keyword search (BM25)

Databases to Know:

DB	Strength	When to Use
Pinecone	Managed, easy	Fastest to production
Qdrant	Open source, Rust-based	Self-hosted, filtering
Weaviate	GraphQL, multimodal	Complex filtering
ChromaDB	Lightweight	Development/prototyping
pgvector	PostgreSQL extension	If already on Postgres
Milvus	High scale	Large-scale deployments

3.10 MLOps

Factor	Details
Importance	★★★★
Interview Frequency	60% of ML Engineer roles
Study Hours	15–20h

Topics:

- Experiment tracking: MLflow, Weights & Biases
- Model registry: MLflow, Vertex AI Model Registry
- Feature stores: Feast, Vertex AI Feature Store, Tecton
- CI/CD for ML: GitHub Actions, Cloud Build, testing ML models
- Model serving: FastAPI, BentoML, Seldon, KServe, Triton Inference Server
- Model monitoring: Evidently AI, WhyLabs, Arize
- Data versioning: DVC
- Orchestration: Airflow, Prefect, Kubeflow Pipelines, Vertex AI Pipelines

3.11 ML System Design

This is the make-or-break round for Senior roles

Framework to answer any ML system design question:



- 1. Problem Scoping (5 min)
 - Clarify requirements, constraints, scale
 - Define success metrics (both ML and business metrics)
- 2. High-Level Architecture (10 min)
 - Data pipeline, feature engineering, model training, serving
 - Draw boxes and arrows
- 3. Data Strategy (5 min)
 - Data sources, labeling, class imbalance
- 4. Model Selection (5 min)
 - Which model family? Why? Tradeoffs?
- 5. Training Infrastructure (5 min)
 - Batch vs online, distributed training
- 6. Serving Infrastructure (5 min)
 - Latency requirements, throughput, caching
- 7. Monitoring & Iteration (5 min)
 - Data drift, model degradation, retraining triggers

Top 10 ML System Design Questions to Practice:

- 1. Design a recommendation system for Netflix/YouTube
- 2. Design a real-time fraud detection system
- 3. Design a semantic search engine for an enterprise document store
- 4. Design a customer support chatbot with RAG
- 5. Design an LLM evaluation pipeline
- 6. Design a multi-tenant AI platform (SaaS)
- 7. Design a content moderation system
- 8. Design a demand forecasting system
- 9. Design a code generation assistant (Copilot-like)
- 10. Design an autonomous AI agent for market research

SECTION 4 — 5 HIGH-IMPACT PROJECTS


PROJECT 1: Enterprise RAG Platform with Evaluation Framework

Target roles: GenAI Engineer, LLM Engineer, AI Engineer

Business Use Case

A production-grade document intelligence platform where enterprises upload PDFs, DOCX, and URLs and query them with natural language. Includes automated evaluation, source citation, and confidence scores.

Architecture



Tech Stack

- **Backend:** FastAPI, Python 3.11+
- **LLM:** Gemini 2.5 Pro via Vertex AI
- **Embeddings:** BGE-large-en-v1.5 (HuggingFace)
- **Vector DB:** Qdrant (self-hosted on Cloud Run)
- **Reranker:** Cohere Rerank v3
- **Evaluation:** RAGAS framework
- **Frontend:** Streamlit or Next.js
- **Deployment:** GCP Cloud Run + Cloud Storage

Resume Bullets



- Built enterprise RAG platform processing 10,000+ documents, achieving 87% faithfulness score on RAGAS evaluation benchmark
- Implemented hybrid retrieval (dense BGE-large + BM25) with Cohere reranking, reducing hallucination rate by 34%
- Deployed on GCP Cloud Run with auto-scaling, serving <2s P95 latency at 500 QPS
- Built automated evaluation pipeline with 4 RAGAS metrics for continuous quality monitoring

Estimated Completion: 2 weekends (16–20h)

PROJECT 2: Multi-Agent Research & Report Generation System

Target roles: GenAI Engineer, AI Agent Engineer, LLM Engineer

Business Use Case

An agentic system that takes a research brief as input and autonomously: searches the web, reads papers, synthesizes information, fact-checks, and produces a structured research report with citations — similar to a Perplexity Pro research workflow.

Architecture (LangGraph)



Tech Stack

- **Orchestration:** LangGraph (StateGraph)
- **LLM:** GPT-4o / Gemini 2.5 Pro
- **Tools:** Tavily Search API, arXiv API, web scraping (Playwright)
- **Memory:** LangGraph checkpointer (PostgreSQL)
- **Backend:** FastAPI with WebSocket for streaming
- **Frontend:** Streamlit with streaming output
- **Deployment:** GCP Cloud Run

Resume Bullets



- Designed 6-agent LangGraph system for autonomous research, reducing research time from 4h to 12 minutes for complex topics
- Implemented supervisor-worker architecture with conditional routing, supporting parallel agent execution
- Built human-in-the-loop checkpointing for fact verification on critical claims
- Reduced hallucinations via multi-agent debate pattern with cross-checking agents

Estimated Completion: 3 weekends (20–25h)


PROJECT 3: AI-Powered Marketing Campaign Automation

Target roles: AI Engineer, Data Scientist, ML Engineer

Business Use Case

End-to-end marketing automation: customer segmentation → audience targeting → personalized content generation → campaign performance prediction → A/B test design.

Architecture



Tech Stack

- **ML:** XGBoost (CTR prediction), scikit-learn (segmentation)
- **GenAI:** Vertex AI (content generation), prompt templates
- **Backend:** FastAPI, Celery for async tasks
- **Data:** BigQuery, Pandas, feature engineering pipeline
- **MLOps:** MLflow for experiment tracking
- **Deployment:** GCP Cloud Run + Cloud Scheduler

Resume Bullets



- Built end-to-end AI marketing platform combining customer segmentation ML models with LLM-powered content personalization
- CTR prediction model (XGBoost) achieved 0.78 AUC vs 0.61 baseline
- LLM content generation pipeline reduced campaign creation time from 3 days to 2 hours
- Implemented MLflow tracking with automated retraining on weekly data refresh

Estimated Completion: 2 weekends (14–18h)


PROJECT 4: Production ML System with Full MLOps Stack

Target roles: ML Engineer, MLOps Engineer, Senior Data Scientist

Business Use Case

A production churn prediction system with full MLOps: feature store, experiment tracking, model registry, automated retraining, monitoring, and drift detection.

Architecture



Tech Stack

- **Features:** Feast feature store
- **Training:** XGBoost + Vertex AI Pipelines (Kubeflow)
- **Tracking:** MLflow
- **Serving:** FastAPI + Docker + GCP Cloud Run
- **Monitoring:** Evidently AI for drift detection
- **Infra:** Terraform for GCP resources
- **CI/CD:** GitHub Actions

Resume Bullets



- Built production ML system with complete MLOps stack: feature store (Feast), experiment tracking (MLflow), automated retraining, and drift monitoring
- Achieved <50ms inference latency with horizontal scaling on Cloud Run
- Implemented data drift detection triggering automated retraining when PSI > 0.2
- CI/CD pipeline with automated model validation tests before production deployment

Estimated Completion: 3 weekends (20–25h)

PROJECT 5: Enterprise AI Knowledge Assistant (Multi-Tenant SaaS)

Target roles: AI Engineer, LLM Engineer, GenAI Engineer, Senior DS

Business Use Case

A multi-tenant AI assistant platform where companies onboard their internal knowledge (Confluence, Notion, Google Drive, Slack) and their employees query it with natural language. Includes access control, usage analytics, and feedback loops.

Architecture



Tech Stack

- **Backend:** FastAPI, PostgreSQL (tenant config), Redis (caching)
- **Vector Store:** Qdrant with multi-collection isolation
- **Auth:** JWT + RBAC
- **LLM:** Vertex AI Gemini 2.5 Pro
- **Analytics:** Posthog or custom Streamlit dashboard
- **Deployment:** GCP Cloud Run, Cloud SQL, GCS
- **Connectors:** Confluence REST API, Google Drive API, Slack API

Resume Bullets



- Built multi-tenant AI knowledge assistant platform serving 3 enterprise clients with data isolation and RBAC
- Implemented connector integrations for Confluence, Google Drive, and Slack with incremental sync
- Designed feedback-driven improvement loop: collected 500+ interactions to improve retrieval quality by 28%
- Production system with 99.5% uptime, <3s P95 latency, serving 200+ daily active users

Estimated Completion: 4 weekends (28–35h)

SECTION 5 — RESUME STRATEGY

5.1 Powerful Resume Summary Examples

For UAE/Dubai Applications:



Senior AI/LLM Engineer with 5+ years building production-grade AI systems at [MNC]. IIT Bombay M.Tech. Specialized in Generative AI, LLM applications, and multi-agent systems. Built RAG pipelines serving 10K+ daily queries, LangGraph-based agentic workflows, and ML systems at scale on GCP/Vertex AI. Seeking senior AI engineering role in Dubai/UAE to drive enterprise AI transformation.

For UK Applications:



Senior Data Scientist / GenAI Engineer | 5 years at [MNC] | IIT Bombay M.Tech
Specialized in production LLM systems, RAG architectures, and multi-agent AI. Experienced across the full ML lifecycle: Python, PyTorch, FastAPI, GCP, Vertex AI, LangGraph. Track record of delivering AI products that reduced operational costs by X% and improved process efficiency by Y%. Open to Skilled Worker Visa sponsorship.

5.2 Achievement-Oriented Bullet Framework

Formula: [Action Verb] + [What You Built/Did] + [Technology] + [Quantified Impact]

Strong Examples:



- ✓ Architected production RAG pipeline using LangGraph + Vertex AI, processing 50K documents and reducing support ticket resolution time by 40%
- ✓ Fine-tuned Llama 3 8B using QLoRA for domain-specific NLP tasks, achieving 12% improvement over GPT-4o baseline at 95% lower inference cost
- ✓ Built multi-agent orchestration system handling 8 concurrent AI workflows, reducing manual analyst work from 20h/week to 3h/week
- ✓ Implemented real-time ML feature pipeline using Kafka + Feast, serving <10ms feature lookup latency for fraud detection model
- ✓ Led ML model migration to GCP Vertex AI, reducing model deployment time from 3 days to 4 hours through automated CI/CD pipeline

5.3 Critical ATS Keywords for AI/GenAI Roles (2026)



Generative AI, Large Language Models, LLM, RAG, Retrieval-Augmented Generation, LangChain, LangGraph, LangSmith, Vector Database, Embeddings, Fine-tuning, LoRA, QLoRA, RLHF, DPO, Prompt Engineering, AI Agents, Multi-Agent Systems, FastAPI, Vertex AI, GCP, PyTorch, Transformers, HuggingFace, Qdrant, Pinecone, Weaviate, MLOps, MLflow, Kubeflow, Model Serving, Docker, Kubernetes, NLP, Named Entity Recognition, Semantic Search, Sentence Transformers, OpenAI API, Gemini, Claude, Anthropic, vLLM, TGI, Model Optimization, Feature Store, Data Pipeline, A/B Testing, Model Monitoring, Drift Detection

5.4 Country-Specific Do's and Don'ts

Factor	UAE	UK	Germany	France
Photo	Optional	✗ Never	✓ Include	✓ Common
DOB	Optional	✗ Never	✓ Include	Optional
Marital Status	Optional	✗ Never	Optional	✗ Not common
Nationality	✓ Include	Not typical	Optional	Not typical
Cover Letter	Optional	Optional	✓ Required	✓ Expected
Length	2 pages	2 pages strict	2 pages + cover	1–2 pages
Visa Note	Not needed	✓ Required	✓ "EU Blue Card eligible"	Optional
References	"Available on request"	"Available on request"	May attach reference letters	Not needed

SECTION 6 — JOB APPLICATION STRATEGY

6.1 Platform Priority Matrix

Platform	Best For	Daily Time	Target	Priority
LinkedIn	UK, Europe, UAE	30 min	3–5 apps	★★★★★
Wellfound (AngelList)	Startups globally	20 min	2–3 apps	★★★★
Indeed	UK, Ireland	15 min	2 apps	★★★
Naukri Gulf	UAE, Gulf	20 min	3 apps	★★★★
GulfTalent	UAE, Gulf	15 min	2 apps	★★★
Bayt.com	UAE, MENA	15 min	2 apps	★★★
Stepstone	Germany	15 min	2 apps	★★★
Xing	Germany	10 min	1 app	★★
Glassdoor	All	10 min	1 app	★★
Company websites	All	30 min	2–3 apps	★★★★★

Golden Rule: Always apply directly on the company website AFTER applying via LinkedIn. Dual application = 2x chance of being seen.

6.2 Best Times to Apply

- **UAE:** Sunday–Tuesday 9–11 AM GST (start of Gulf work week)
 - **UK:** Tuesday–Thursday 8–10 AM BST
 - **Germany:** Monday–Wednesday 9–11 AM CET
 - **LinkedIn Easy Apply:** Tuesday–Thursday mornings get highest recruiter inbox attention
-

6.3 Outreach Templates

Cold Outreach to Hiring Manager (LinkedIn)



Hi [Name],

I noticed [Company] is expanding its AI/GenAI capabilities — [specific news/job post].

I'm a Senior AI Engineer with 5 years at [MNC], IIT Bombay M.Tech, building production LLM systems (RAG pipelines, multi-agent systems with LangGraph, GCP/Vertex AI).

I've specifically admired [specific thing about their AI work/product].

Would you be open to a 15-minute chat? I'd love to understand more about the team's direction. Happy to share my portfolio.

Best, [Name]

Referral Request Template



Hi [Name],

Hope you're doing well! I noticed you're at [Company] — I've been following their AI work closely, especially [specific product/initiative].

I'm actively looking for GenAI Engineer / Senior Data Scientist roles in [Dubai/London/etc.]. Given my 5 years building LLM systems at [MNC] (IIT Bombay M.Tech), I think I'd be a strong fit.

Would you be comfortable referring me, or pointing me to the right person to speak with? I can share my resume/portfolio.

No pressure at all if it doesn't feel right — just thought it was worth asking!

Thanks,
[Name]

Follow-up After No Response (2 weeks later)



Hi [Name],

Just a gentle follow-up on my message from [date]. I know you're busy!

Quick summary: Senior AI/LLM Engineer, 5 YOE, IIT Bombay M.Tech, interested in [specific role/team].

I've since [deployed RAG project / completed XYZ] which might be relevant to what your team is building.

Happy to connect for 10 minutes. If timing isn't right, no worries — appreciate you!

Best,
[Name]

6.4 How to Identify Scam/Fake Jobs

- Red Flags:
- Salary "too good to be true" (e.g., AED 100K/month for junior role)
 - No company name listed, only "Confidential Company"
 - Grammar errors in job description
 - Asks for personal documents before any interview
 - Contact is personal Gmail/Yahoo (not corporate email)
 - Job posted multiple times with different company names
 - Asks you to pay for visa/background check upfront
- Verification Steps:
1. Search company name on LinkedIn — does it exist?
 2. Cross-reference job on company's official website

- 3. Check if recruiter's LinkedIn shows they work at that company
- 4. Never pay for visa processing before receiving an official contract

6.5 Response Rate Optimization

Target: 5–10% response rate (industry average is 2–3%)

Tactics to increase above average:

- 1. Tailor top 1/3 of resume to each job description — 3 minutes of tailoring = 3x response rate
- 2. Apply within 24 hours of posting — early applicants get 3–4x more views
- 3. Send connection request to HR/recruiter at that company same day as applying
- 4. Reference specific company project/product in cover note
- 5. Add "Open to Work" recruiter badge (not public — only recruiters see it)
- 6. Post on LinkedIn the same week — active profiles get more recruiter views

SECTION 7 — INTERVIEW MASTER PLAN

7.1 Typical Interview Flows by Country

UAE Interview Flow (3–5 weeks total)



- Day 0: Apply
- Day 3–7: Recruiter screening call (30 min)
- Day 7–14: HR call (culture, salary, visa, notice period)
- Day 14–21: Technical round 1 (ML/Python/GenAI concepts, 60 min)
- Day 21–28: Take-home assignment OR technical round 2 (system design)
- Day 28–35: Final round (leadership/director/hiring manager)
- Day 35–42: Offer

UK Interview Flow (6–12 weeks total)



- Day 0: Apply
- Day 7–14: Recruiter screening (phone)
- Day 14–21: Technical screening (coding + ML, 45–60 min)
- Day 21–35: Technical round (system design + deep dive, 90 min)
- Day 35–49: Final panel (HR + Tech Lead + Manager, 3–4 hours)
- Day 49–70: Offer + visa process starts

Germany Interview Flow (8–16 weeks total)



- Day 0: Apply + cover letter
- Day 14–21: HR intro call
- Day 21–35: Technical assessment (take-home OR online coding)
- Day 35–49: Technical interview (video call, 60–90 min)
- Day 49–63: Onsite/video panel (3–4 interviews on one day)
- Day 63–84: Offer + EU Blue Card process (4–6 weeks for visa)

7.2 Technical Interview Question Bank (2026)

LLM/GenAI Questions (Most Frequent)

- "Explain transformer architecture. What is multi-head attention?"
- "What is RAG? When would you prefer RAG over fine-tuning?"
- "How do you evaluate a RAG system's quality?"
- "What is LangGraph? How does it differ from LangChain?"
- "How would you prevent prompt injection in a production AI system?"
- "Explain the difference between LoRA and QLoRA. When would you use each?"
- "How does KV cache work? Why is it important for inference efficiency?"
- "What is DPO? How does it compare to RLHF with PPO?"
- "Design a multi-agent system for customer support. What agents would you have?"
- "How do you handle hallucinations in a production RAG system?"

ML Fundamentals Questions

- "Explain bias-variance tradeoff with a concrete example"
- "How does XGBoost differ from Random Forest? When would you choose each?"
- "Walk me through gradient descent optimization. What are the variants?"
- "How would you handle severely imbalanced classes (1:1000 ratio)?"
- "Explain SHAP values and how you'd use them for model explainability"
- "Design an A/B test for a new recommendation model"

System Design Questions

- "Design a real-time recommendation engine for 10M users"
- "Design an LLM-powered document search system for an enterprise"
- "Design a fraud detection system with sub-100ms latency requirement"
- "Design a multi-tenant AI platform where each customer has isolated data"
- "Design a model monitoring system that detects concept drift"

7.3 Salary Negotiation Strategy

The Golden Rules:

- Never give a number first** — "I'd like to understand the full package before discussing a specific number"
- Always anchor high** — Research P90 salary for the role. Quote 10–15% above your target

3. **Use competing offers** as leverage (even if UAE vs UK — they're different markets)
4. **Negotiate the full package:** base + equity/ESOP + relocation allowance + visa costs + signing bonus

UAE Negotiation Script:



"Based on my research and the scope of this role, I'm targeting AED 45,000–50,000/month. I'm also considering opportunities in the UK at £90,000–£100,000, but I'm particularly excited about this role's focus on [specific thing]. I'm confident we can find a number that works for both sides."

Visa Sponsorship Discussion:



"I want to be upfront — I'll require visa sponsorship. Based on my research, the [UAE Employment Visa / UK Skilled Worker Visa] process typically takes [2–4 / 6–8] weeks and is straightforward for tech roles at your level. I've already verified I meet all eligibility requirements. Is this something your team handles routinely?"

SECTION 8 — DAILY STUDY SCHEDULE

8.1 Balanced Version (Working Professional)

Target: 4–5 hours/weekday, 8–9 hours/weekend



WEEKDAY SCHEDULE:
05:30 – Wake up, no screens for 20 min
06:00–08:00 – Deep technical study (LLM/ML/coding) [2h] ← MOST IMPORTANT
08:00–09:00 – Commute/breakfast
09:00–18:00 – Day job
18:30–19:30 – Interview prep or project work [1h]
19:30–20:30 – Job applications + outreach [1h]
20:30–21:30 – LeetCode OR resume optimization [1h]
21:30 – Wind down, review notes

WEEKEND SCHEDULE:
09:00–12:00 – Project building [3h]
13:00–17:00 – Intensive technical study [4h]
17:00–19:00 – Job applications (5 per day) [2h]
19:00–21:00 – Mock interview practice OR LinkedIn content [2h]

8.2 Aggressive Version (If Job Searching Full-Time)

Target: 8–10 hours/day



07:00–09:00 – LeetCode + Python/ML fundamentals [2h]
09:00–12:00 – Project building [3h]
13:00–15:00 – LLM/GenAI deep study [2h]
15:00–17:00 – Job applications (8–10) [2h]
17:00–18:00 – Recruiter outreach + LinkedIn [1h]
19:00–21:00 – System design practice [2h]

8.3 Burnout Prevention

- **One full day off per week** (no studying, no applications)
- **Exercise daily** — 30 min minimum. Non-negotiable for cognitive performance
- **No studying after 10 PM** — quality degrades
- **Pomodoro technique**: 25 min focused work + 5 min break
- **Weekly retrospective**: What worked? What didn't? Adjust plan
- **Celebrate small wins**: First recruiter call, first interview, first offer = reward yourself

SECTION 9 — TRACKERS & TEMPLATES

9.1 Job Application Tracker

Excel Sheet: Tab "Applications"

Column	Values
Date Applied	DD/MM/YYYY
Company	Name
Role Title	Exact JD title
Location	Dubai / London / etc.
Platform	LinkedIn / Direct / Naukri Gulf
JD URL	Link
Resume Version Used	v1.0 / v1.1-UAE / v1.2-UK
Application Status	Applied / Screening / Tech Round / Offer / Rejected
Recruiter Name	Name
Recruiter Email/LinkedIn	Contact
Follow-up Date	2 weeks after apply
Notes	Key JD keywords, interview feedback
Priority	High / Medium / Low
Salary Range	From JD or research

Color Coding:

-  Green = Active / Positive progress
 -  Yellow = Applied / Awaiting response
 -  Red = Rejected / No response after 3 weeks
 -  Blue = On hold / Future consideration
-

9.2 Recruiter Outreach Tracker

Excel Sheet: Tab "Outreach"

Column	Values
Date	DD/MM/YYYY
Name	Recruiter/HM name
Company	
LinkedIn URL	
Platform	LinkedIn / Email
Message Sent	Yes/No
Response	Yes/No/Not Yet
Response Date	
Outcome	Call Scheduled / Referred / No Fit / Pending
Follow-up #1 Date +14 days	
Follow-up #2 Date +28 days	
Notes	

9.3 Interview Tracker

Excel Sheet: Tab "Interviews"

Column	Values
Company	
Round #	1, 2, 3...
Interview Type	HR / Technical / System Design / Case Study
Date & Time	
Interviewer Name	
Format	Video / Phone / In-person
Status	Scheduled / Completed / Passed / Failed
Questions Asked	Detailed notes
My Answers	What I said
Feedback	If received
Next Steps	
Offer?	Yes/No/Pending

9.4 Daily Study Tracker

Excel Sheet: Tab "Daily Progress"

Column	Values
Date	
Topics Studied	
Study Hours	
LeetCode Solved	# and difficulty
Applications Sent	#
Outreach Messages	#
Project Progress	% complete
LinkedIn Activity	Post / Connection / Comment
Mood / Energy	1–5
Tomorrow's Priority	Top 3 tasks

9.5 Weekly Dashboard

Excel Sheet: Tab "Weekly KPIs"

KPI	Week 1	Week 2	Week 3	Week 4	Target/Week
Applications Sent					15
Response Rate %					>5%
Recruiter Screens					2
Technical Rounds					1
Offers					0 → 1 by W8
Study Hours					30h
LeetCode Problems					10
LinkedIn Connections					20
LinkedIn Post Views					500
GitHub Commits					20

SECTION 10 — FINAL EXECUTION BLUEPRINT

10.1 Top 20 Most Important Actions (Priority Order)

1. **Update LinkedIn headline RIGHT NOW** — "Senior AI/LLM Engineer | Open to Dubai/UK/Europe"
2. **Create ATS-optimized resume** with quantified bullets before first application
3. **Build Project 1 (Enterprise RAG)** — single most impressive portfolio piece for 2026
4. **Connect with 20 UAE tech recruiters this week** via LinkedIn DM
5. **Apply to 5 roles per day minimum** starting Week 2
6. **Master LangGraph** — it's the differentiator skill in 2026 agent roles
7. **Practice ML system design daily** — it's the make-or-break round for senior roles
8. **Write LinkedIn posts weekly** — active technical posts get you inbound
9. **Set up application tracker** — manage 100+ applications requires a system
10. **Solve 5 LeetCode problems per week** — companies still test coding
11. **Research each company before applying** — tailored applications convert 3x better
12. **Never state expected salary first** — always get their range first
13. **Be explicit about visa needs** upfront with recruiters — don't waste 5 rounds then get rejected on visa
14. **Build Project 2 (Multi-Agent)** — shows you're ahead of the market
15. **Follow up every 2 weeks** on applications — 30% of responses come from follow-ups

- 16. Practice "tell me about yourself" 20 times until it's natural
- 17. Know the company's AI strategy before every interview
- 18. Prepare 10 STAR behavioral stories covering leadership, failure, impact
- 19. Deploy every project on GCP — cloud deployment = production credibility
- 20. Take care of your body — exercise, sleep, nutrition. The 60-day sprint is a marathon

10.2 Biggest Mistakes to Avoid

Mistake	Impact	Correction
Applying without tailoring resume	10x lower response rate	Spend 5 min tailoring top section
Not following up after applying	Miss 30% opportunities	Set calendar reminder for 2 weeks
Being vague about visa status	Wasted interview cycles	State upfront: "I require sponsorship"
Skimming company research	Failed final rounds	Spend 30 min on each company before interview
Underestimating UAE salary range	Leave money on table	Research AED 35–55K range for your level
Only applying to big tech	Low conversion	Balance startups + consulting + product companies
Practicing theory, not coding	Fail technical screen	Solve 2 LeetCode problems DAILY
Not building projects	Weak portfolio	Projects = proof of skills
Waiting for perfect resume	Delay in applications	Good resume now > perfect resume in 2 months
Neglecting LinkedIn activity	Low recruiter inbound	Post weekly, comment daily

10.3 Realistic Timeline to Outcomes

Outcome	Realistic Timeline	Optimistic
First recruiter contact	Week 2–3	Week 1–2
First recruiter screen	Week 3–4	Week 2–3
First technical interview	Week 4–6	Week 3–5
First UAE offer	Week 6–10	Week 5–8
First UK offer	Week 8–14	Week 7–12
First Germany offer	Week 10–18	Week 9–16
Multiple competing offers	Week 10–16	Week 8–14

10.4 Success Probability Analysis

Country	Probability of Getting Offer in 90 Days	Key Variables
UAE/Dubai	75–85%	Very favorable for Indians, fast process
UK	55–70%	Good, requires visa sponsor. IIT + MNC = strong
Netherlands	50–65%	Fast IND process, 30% ruling attractive
Ireland	55–70%	Big tech EU hubs, Critical Skills permit
Germany	40–55%	Slower, language can be barrier
France	35–50%	Talent Passport good but French language matters

Overall: 85%+ probability of at least 1 international offer within 90 days given consistent execution

10.5 Country ROI Ranking (Take-Home vs Effort)

1. 🇦🇪 **UAE/Dubai** — Tax-free, fastest hiring, largest Indian community, easiest visa
2. 🇳🇱 **Netherlands** — 30% tax ruling, fast IND visa, English-only tech
3. 🇮🇪 **Ireland** — Big tech EU HQs, Critical Skills permit, English-only
4. 🇬🇧 **UK** — Large market, great GenAI ecosystem, but higher competition
5. 🇩🇪 **Germany** — Lower net salary after tax, longer process, language matters

10.6 Backup Strategies

If no offers by Day 45:

- Increase application volume by 50%
- Expand to more countries (Canada, Singapore, Australia)
- Consider remote-first international companies (no relocation required)
- Reach out to Indian alumni at target companies (IIT Bombay alumni network is gold)
- Consider contract/freelance roles as entry point (especially UAE)

If no technical interviews by Day 30:

- Resume is likely the issue → Get it reviewed by 3 people + ATS checker
- LinkedIn needs optimization → Audit for keywords
- Application quality issue → More tailoring, less spray-and-pray

IIT Bombay Alumni Network (Critical Asset):

- iitbaa.com — official alumni network
- LinkedIn search: "IIT Bombay" + "[target company]"
- Alumni are statistically far more likely to refer you. Target 10 outreaches to IIT Bombay alumni per week

10.7 Final Action Items — Start TODAY



- TODAY (May 10, 2026):
- ❑ Update LinkedIn headline to include "Open to Dubai/UK/Europe"
 - ❑ Download ATS-checker tool (jobscan.co — free version)
 - ❑ Create Excel tracker with all tabs from Section 9
 - ❑ Set up LinkedIn job alerts for: "GenAI Engineer", "LLM Engineer", "Senior Data Scientist" in Dubai, London, Amsterdam, Berlin, Dublin
 - ❑ Write list of 50 target companies across all target countries
 - ❑ Start writing resume v1.0

- THIS WEEK (May 10–16):
- ❑ Complete LinkedIn to All-Star status
 - ❑ Resume v1.0 and v1.1-UAE ready
 - ❑ GitHub cleaned and 3 repos pinned
 - ❑ 20 recruiters connected on LinkedIn
 - ❑ 5 applications submitted
 - ❑ Project 1 architecture designed

- BY DAY 30 (June 9):
- ❑ 50+ applications submitted
 - ❑ Project 1 and 2 deployed on GitHub + GCP
 - ❑ 3 recruiter screens completed
 - ❑ First technical interview done
 - ❑ 40+ LeetCode problems solved
 - ❑ 5 LinkedIn posts published

- BY DAY 60 (July 9):
- ❑ 115+ applications submitted
 - ❑ 5 projects deployed
 - ❑ 10+ interview rounds completed
 - ❑ 1–3 offers in hand or very close
 - ❑ Negotiating salary and visa terms

"The best time to start was yesterday. The second best time is NOW."

Your IIT Bombay background + MNC experience + GenAI skills is a top-3% profile globally for these roles. The only variable is execution consistency. Follow this plan and you WILL have international offers in 60–90 days.

Document Version: 1.0 | **Created:** May 2026 | **Review Every:** Sunday